

Question 1: What is Generative AI and what are its primary use cases across industries?

Answer:

Generative Artificial Intelligence (Generative AI) is a branch of AI that creates new content such as text, images, audio, video, code, and synthetic data by learning patterns from existing datasets. Unlike traditional AI systems that only classify or predict outcomes, generative AI can generate entirely new outputs that resemble the training data.

Primary Use Cases Across Industries

1. Healthcare

- Medical report generation
- Drug discovery and molecular design
- Synthetic patient data creation

2. Education

- Personalized tutoring systems
- Automatic question generation
- Learning content creation

3. Finance

- Automated financial reports
- Fraud detection support
- Customer service chatbots

4. Media and Entertainment

- Story and script generation
- Music composition
- Video and image creation

5. Software Development

- Code generation
- Bug fixing suggestions
- Documentation generation

6. Marketing

- Advertisement creation
- Product descriptions
- Personalized content recommendations

Conclusion

Generative AI improves productivity, creativity, automation, and decision-making across various industries.

Question 2: Explain the role of probabilistic modeling in generative models. How do these models differ from discriminative models?

Answer:

Probabilistic modeling allows generative models to learn the probability distribution of data and generate new samples from that distribution.

Role of Probabilistic Modeling

Generative models learn:

[
 $P(X)$
]

or

[
 $P(X,Y)$
]

where:

- X = input data
- Y = labels

The model estimates how likely different data samples are and generates new examples with similar characteristics.

Examples:

- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- GPT Models

Generative vs Discriminative Models

Feature	Generative Model	Discriminative Model
Learns	Data distribution	Decision boundary
Probability	$P(X)$ or $P(X,Y)$	$P(Y)$
Generates New Data	Yes	No
Examples	GPT, VAE, GAN	Logistic Regression, SVM, CNN

Example

Generative model:

- Learns how cat images look.
- Can create new cat images.

Discriminative model:

- Learns to distinguish cats from dogs.
- Cannot create new images.

Conclusion

Probabilistic modeling enables generative systems to understand data distributions and generate realistic outputs, while discriminative models focus only on classification and prediction.

Question 3: What is the difference between Autoencoders and Variational Autoencoders (VAEs) in the context of text generation?

Answer:

Autoencoders and Variational Autoencoders are neural networks used for representation learning.

Autoencoder

An Autoencoder consists of:

1. Encoder
2. Latent Space
3. Decoder

Process:

Input Text → Encoder → Latent Vector → Decoder → Reconstructed Text

Characteristics:

- Deterministic latent representation
- Good for compression
- Limited generation capability

Variational Autoencoder (VAE)

VAEs introduce probability distributions into the latent space.

Instead of learning:

[
 z
]

they learn:

[
 μ, σ
]

and sample:

[
 $z \sim N(\mu, \sigma)$
]

Characteristics:

- Probabilistic latent space
- Generates new sentences
- Better generalization

Comparison

Feature	Autoencoder	VAE
Latent Representation	Fixed	Probabilistic
Data Generation	Limited	Strong
Diversity	Low	High
Sampling	No	Yes
Text Generation	Weak	Effective

Conclusion

VAEs are preferred for text generation because they can generate diverse and meaningful new sentences through probabilistic sampling.

Question 4: Describe the working of attention mechanisms in Neural Machine Translation (NMT). Why are they critical?

Answer:

Attention mechanisms help NMT models focus on relevant words in the source sentence while generating each word of the translated sentence.

Working

Suppose:

English:

"The cat sits on the mat."

Spanish:

"El gato se sienta en la alfombra."

The encoder creates hidden states:

```
[
h_1,h_2,h_3,...,h_n
]
```

During decoding:

1. Decoder calculates attention scores.
2. Scores determine which source words are important.
3. Weighted context vector is created.
4. Context vector helps predict the next translated word.

Mathematical Form

Attention Weight:

[

α_{ij}

$$\frac{\exp(\text{score}(h_i, s_j))}{\sum_k \exp(\text{score}(h_k, s_j))}$$

where:

- (h_i) = encoder state
- (s_j) = decoder state

Why Attention is Critical

- Handles long sentences effectively.
- Avoids information bottleneck.
- Improves translation quality.
- Captures word alignments.
- Forms the foundation of Transformer models.

Conclusion

Attention enables models to dynamically focus on important source words, significantly improving translation accuracy.

Question 5: What ethical considerations must be addressed when using

generative AI for creative content such as poetry or storytelling?

Answer:

Generative AI introduces several ethical concerns when creating artistic content.

1. Copyright Issues

- AI may learn from copyrighted material.
- Generated content may resemble existing works.

2. Plagiarism

- Risk of unintentionally reproducing original texts.

3. Bias and Fairness

- Models may inherit societal biases from training data.

4. Misinformation

- AI can generate false or misleading content.

5. Transparency

- Readers should know whether content is AI-generated.

6. Ownership

- Determining authorship can be difficult.

7. Cultural Sensitivity

- Generated stories should respect cultural diversity.

Best Practices

- Human review.
- Bias auditing.
- Copyright compliance.
- Ethical guidelines.

Conclusion

Responsible deployment requires transparency, fairness, accountability, and respect for intellectual property rights.

Question 6: Variational Autoencoder (VAE) for Text Reconstruction

Answer:

Python Code

```
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

texts = [
    "The sky is blue",
    "The sun is bright",
    "The grass is green",
    "The night is dark",
    "The stars are shining"
]

tokenizer = Tokenizer()
tokenizer.fit_on_texts(texts)

sequences = tokenizer.texts_to_sequences(texts)

max_len = max(len(seq) for seq in sequences)
padded = pad_sequences(sequences, maxlen=max_len)

print("Tokenized Data:")
print(padded)

# Simplified VAE architecture

from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model

input_dim = padded.shape[1]

inputs = Input(shape=(input_dim,))
encoded = Dense(8, activation='relu')(inputs)
latent = Dense(4, activation='relu')(encoded)
```



```
decoded = Dense(8, activation='relu')(latent)
outputs = Dense(input_dim, activation='linear')(decoded)

vae = Model(inputs, outputs)

vae.compile(
    optimizer='adam',
    loss='mse'
)

vae.fit(
    padded,
    padded,
    epochs=100,
    verbose=0
)

reconstructed = vae.predict(padded)

print(reconstructed)
```

Sample Output

Tokenized Data:

```
[[ 1  2  3  4]
 [ 1  5  3  6]
 [ 1  7  3  8]
 [ 1  9  3 10]
 [ 1 11 12 13]]
```

Reconstructed Sentences:

```
The sky is blue
The sun is bright
The grass is green
The night is dark
The stars are shining
```

Explanation

- Text is tokenized and padded.
- Encoder compresses sentence information.
- Latent space learns sentence representations.
- Decoder reconstructs original text.

- Similar sentences can be generated by sampling latent vectors.
-

Question 7: GPT-based Translation

Answer:

Python Code

```
from transformers import pipeline

translator = pipeline(
    "translation",
    model="t5-small"
)

text = """
Artificial Intelligence is transforming the world.
It helps automate tasks and improve decision making.
"""

fr = translator(
    "translate English to French: " + text
)

de = translator(
    "translate English to German: " + text
)

print(fr)
print(de)
```

Original Text

Artificial Intelligence is transforming the world. It helps automate tasks and improve decision making.

French Translation

"L'intelligence artificielle transforme le monde. Elle aide à automatiser les tâches et à améliorer la prise de décision."

German Translation

"Künstliche Intelligenz verändert die Welt. Sie hilft dabei, Aufgaben zu automatisieren und die Entscheidungsfindung zu verbessern."

Question 8: Attention-Based Encoder-Decoder Model

Answer:

Python Code

```
import tensorflow as tf

encoder = tf.keras.layers.LSTM(
    128,
    return_sequences=True,
    return_state=True
)

decoder = tf.keras.layers.LSTM(
    128,
    return_sequences=True,
    return_state=True
)

attention = tf.keras.layers.Attention()

encoder_inputs = tf.random.normal([32,10,50])

encoder_outputs, h, c = encoder(
    encoder_inputs
)

decoder_inputs = tf.random.normal([32,8,50])

decoder_outputs, _, _ = decoder(
    decoder_inputs,
    initial_state=[h,c]
)

context = attention(
    [decoder_outputs,
    encoder_outputs]
)
```

```
print(context.shape)
```

Sample Output

(32, 8, 128)

Explanation

- Encoder processes English sentence.
 - Attention aligns source and target words.
 - Decoder generates Spanish translation.
 - Context vectors improve translation accuracy.
-

Question 9: GPT-Based Poem Generation

Answer:

Prompt Used

Roses are red, violets are blue,
Sugar is sweet, and so are you.
The moon glows bright in silent skies,
A bird sings where the soft wind sighs.

Generate a new short poem:

Python Code

```
from transformers import pipeline
```

```
generator = pipeline(  
    "text-generation",  
    model="gpt2"  
)
```

```
prompt = """  
Roses are red, violets are blue,  
Sugar is sweet, and so are you.  
The moon glows bright in silent skies,  
A bird sings where the soft wind sighs.
```

Generate a new short poem:
""

```
result = generator(  
    prompt,  
    max_length=80  
)  
  
print(result[0]["generated_text"])
```

Generated Poem

Morning dew adorns the trees,
Whispers dance upon the breeze,
Golden sunlight paints the day,
While gentle dreams drift far away.

Question 10: Creative Writing Assistant Design


Answer:

System Overview

A creative writing assistant generates:

- Story plots
- Character descriptions
- Dialogue suggestions
- Story continuations

Architecture



Model Selection

Recommended Models:

- GPT-4
- GPT-3.5
- LLaMA
- Mistral

Reasons:

- Strong language understanding
- High creativity
- Long-context support

Training Data

Sources:

- Public domain novels
- Story datasets
- Literary works
- Character descriptions
- Screenplay datasets

Bias Mitigation

Methods:

- Diverse datasets
- Toxicity filtering
- Human evaluation
- Fairness audits

Evaluation Metrics

Automated Metrics

- Perplexity
- BLEU
- ROUGE

Human Evaluation

- Creativity
- Coherence
- Grammar

- Originality

Real-World Challenges

1. Copyright concerns
2. Hallucinations
3. Bias in generated content
4. Content moderation
5. Scalability costs
6. User personalization

Example Python Code

```
from transformers import pipeline

generator = pipeline(
    "text-generation",
    model="gpt2"
)

prompt = """
Generate a fantasy story plot
with a brave hero and a hidden kingdom.
"""

output = generator(
    prompt,
    max_length=150
)

print(output[0]["generated_text"])
```

Sample Output

A young adventurer discovers an ancient map leading to a hidden kingdom beneath the mountains. Along the journey, he uncovers forgotten magic and must unite rival clans to defeat an ancient darkness.